

STIR and Tensorflow

Philipp Windischhofer

July 27, 2017

The big picture

- ▶ STIR
- ▶ iterative reconstruction algorithms
- ▶ find the optimum of a *cost function*

$$P(\text{global LOR response}|\text{image}) = \prod_{\text{all LORs}} P(\text{LOR response}|\text{image})$$

$$P(\text{LOR response}|\text{image}) = \int_{\text{LOR}} \text{image density}$$

- ▶ Image (estimate) \rightsquigarrow LOR response: *forward projection*

Ray tracing

- ▶ in the following: how to compute / approximate this integral?
- ▶ represent *image* as a discretized array of voxels

$$\int_{\text{LOR}} \text{image density} \rightsquigarrow \sum_{\text{traversed voxels}} \text{LOI} \cdot \text{voxel brightness}$$

- ▶ $\text{LOI} = \frac{\text{length of intersection}}{\text{voxel}}$ of the LOR through a specific voxel

(Standard) STIR spends $> 90\%$ of the reconstruction time for *forward projection*

- ▶ is there a way to make it faster?

Table of Contents

The big picture

Tensorflow - what is it?


How to do ray tracing on a CPU?

How to do ray tracing on a GPU?

Results

Tensorflow

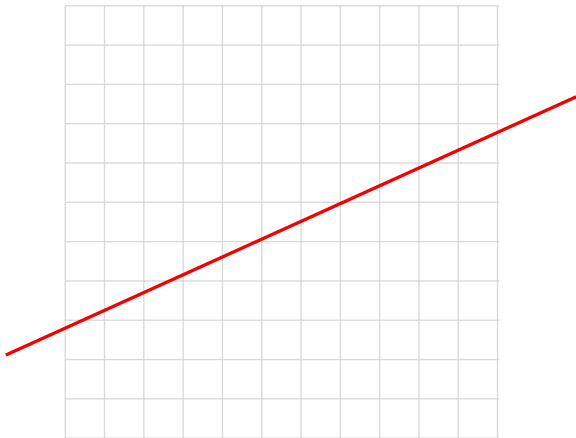
- ▶ *Google*: “An open-source software library for machine intelligence”
 - ▶ heavily used in the machine-learning community
- ▶ Much more than that:
 - ▶ a software package for numerical calculations using data-flow *graphs*



- ▶ Design the graph *now*, save it to a file, run it *later*...
- ▶ Can run it (almost) everywhere: CPU, GPU, a cluster, ...

How to do ray tracing on a CPU?

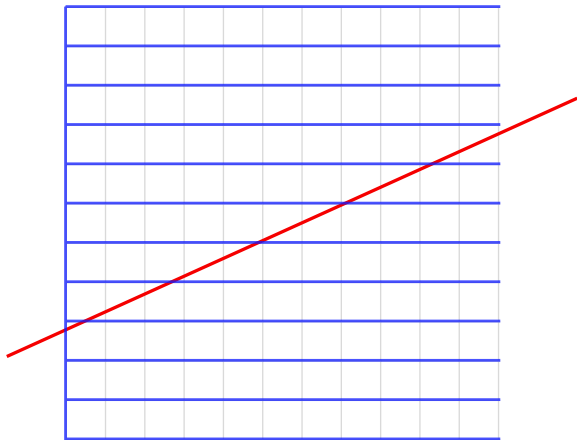
Siddon's algorithm (2D)



- ▶ a (very) sequential algorithm

How to do ray tracing on a CPU?


Siddon's algorithm (2D)



- get intersections of LOI with all horizontal planes

How to do ray tracing on a CPU?

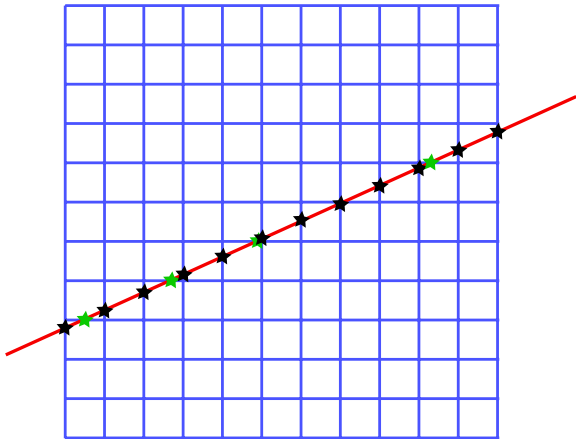
Siddon's algorithm (2D)



- get intersections of LOI with all vertical planes

How to do ray tracing on a CPU?

Siddon's algorithm (2D)



- combine them and get the individual LOIs

How to do ray tracing on a CPU?

Siddon's algorithm (2D)

- ▶ parametrize the LOR

$$\mathbf{x} = (\mathbf{x}_{\text{end}} - \mathbf{x}_{\text{start}}) \cdot \alpha + \mathbf{x}_{\text{start}}$$

$$\alpha \in [0, 1]$$

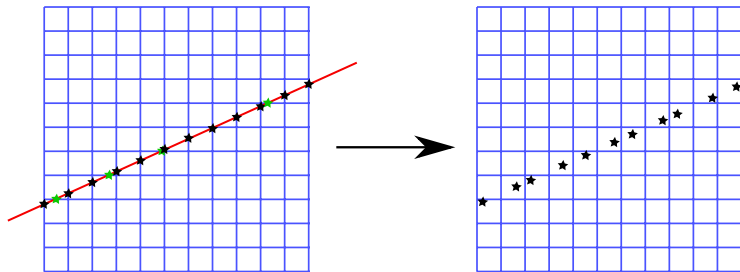
- ▶ get intersections with horizontal and vertical planes in terms of the line parameter α
- ▶ merge and sort them
- ▶ *differences* between neighbouring α values in the merged list describe the LOIs!

Siddon and Tensorflow

- ▶ Siddon: $\mathcal{O}(N)$ for a LOR of length N (this means a voxel array of size N^3 !!)
- ▶ many branches, a loop over the LOR, ...
- ▶ very inefficient when running with Tensorflow (GPU \neq CPU architecture)
 - ▶ especially problematic when trying to trace multiple LORs “in parallel”
- ▶ for Tensorflow: need an algorithm that ...
 - ▶ uses the same algorithmic steps on *all* voxels of the LOR
 - ▶ uses the same algorithmic steps on *all* LORs

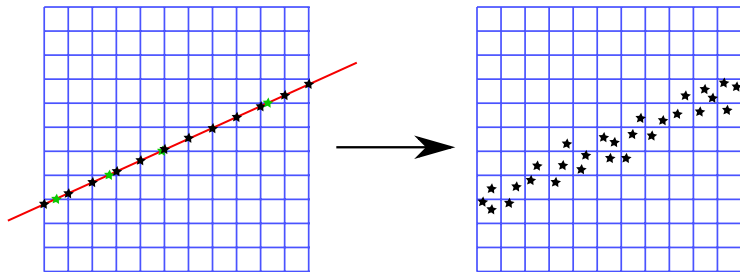
Alternatives to Siddon

- ▶ naive approach: for every LOR, compute the LOI through every voxel, then add them all up
 - ▶ works fine with Tensorflow, but scales as $\mathcal{O}(N^3)$
 - ▶ CPU & Siddon beat it again for real-world array sizes
- ▶ alternative approach: think of (intersection) points *without* the notion of an LOR



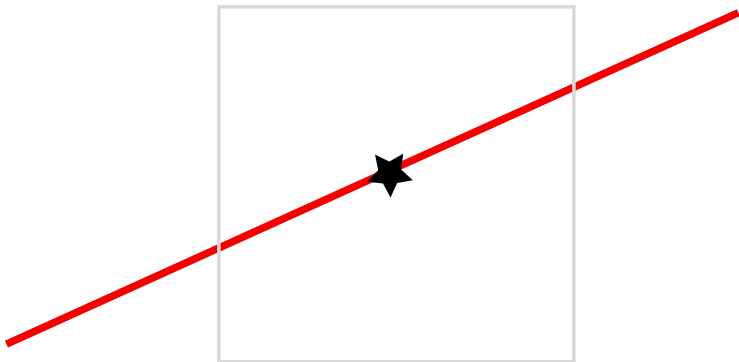
Alternatives to Siddon

- ▶ naive approach: for every LOR, compute the LOI through every voxel, then add them all up
 - ▶ works fine with Tensorflow, but scales as $\mathcal{O}(N^3)$
 - ▶ CPU & Siddon beat it again for real-world array sizes
- ▶ alternative approach: think of (intersection) points *without* the notion of an LOR



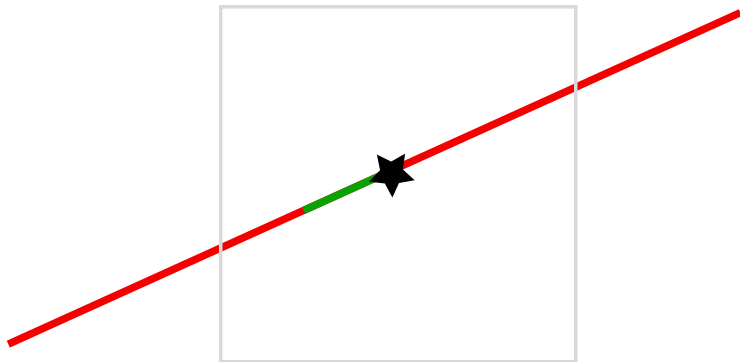
Voxel-by-voxel LOI calculation

- ▶ still not much easier than the full problem!
- ▶ a cube is very discontinuous, hard to do it analytically without branches
- ▶ what if content with an *approximate* solution?



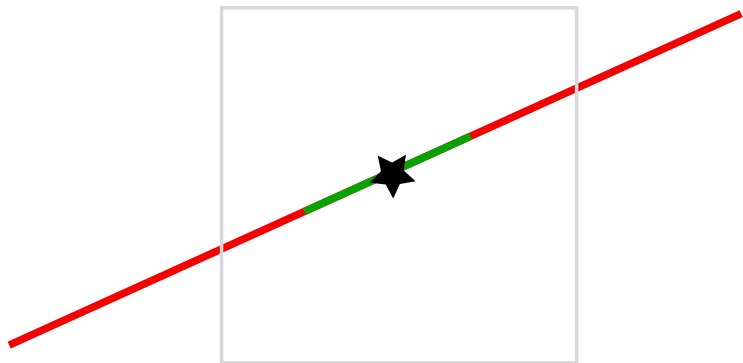
Voxel-by-voxel LOI calculation

- ▶ still not much easier than the full problem!
- ▶ a cube is very discontinuous, hard to do it analytically without branches
- ▶ what if content with an *approximate* solution?




Voxel-by-voxel LOI calculation

- ▶ still not much easier than the full problem!
- ▶ a cube is very discontinuous, hard to do it analytically without branches
- ▶ what if content with an *approximate* solution?



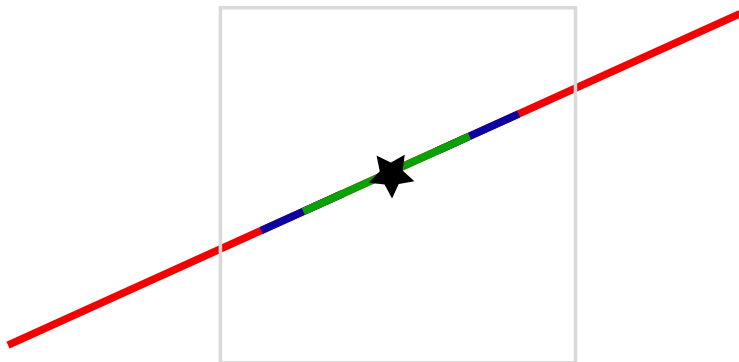
Voxel-by-voxel LOI calculation

- ▶ still not much easier than the full problem!
- ▶ a cube is very discontinuous, hard to do it analytically without branches
- ▶ what if content with an *approximate* solution?



Voxel-by-voxel LOI calculation

- ▶ still not much easier than the full problem!
- ▶ a cube is very discontinuous, hard to do it analytically without branches
- ▶ what if content with an *approximate* solution?




Ray marching: an iterative algorithm

- ▶ Use *signed distance function* of a cube to cut down the number of iterations that are needed
 - ▶ a heuristic that produces an *underestimation* of the distance to the voxel boundary
- ▶ need to assess the residual error that remains

Ray marching: an iterative algorithm

Accuracy

- assume: have “enough” points along the LOR




- with 6 iterations, are already at $\sim 5\%$ level!

Bringing together STIR and Tensorflow

The infrastructure

- ▶ Build the Tensorflow graph in a Python script
 - ▶ much easier and faster to prototype
 - ▶ Python API much better documented and more stable
- ▶ Save the complete graph into a ProtoBuf file
- ▶ Load the file using the C++ API and make use of it in STIR
 - ▶ link STIR against the Tensorflow (shared) library


The workflow



Caching in STIR


- ▶ scanners have (many) symmetries
- ▶ ProjMatrixByBin maintains a cache
 - ▶ already computed matrix elements are stored
 - ▶ new matrix elements related by symmetries to known ones “come for free”!
- ▶ but: caching most efficient if *every* new matrix element request is checked against cache before computing
- ▶ *batching* the matrix element requests makes caching very inefficient!

Speedup without caching



left: 6 iterations, **right:** 2 iterations


- average speedup is similar



left: 6 iterations, **right:** 2 iterations

- ▶ *I/O*: converting from ProjMatrixElemsForOneBin to Tensor and back
- ▶ *point scheduling*: choose points to sample the TOR / LOR

Images



left: 2 iterations, 20LORs per matrix element, **right:** original STIR

How to proceed?

- ▶ demonstrated a proof-of-concept of integrating Tensorflow into STIR
- ▶ whole toolchain is in place
 - ▶ separation of graph generation (Python) and usage (C++) works very well
- ▶ Speedup of ~ 2 observed

Not too bad, but still some limitations:

- ▶ how to improve utilization of GPU (conversion between Tensors and STIR objects)?
- ▶ is there a better way to estimate the matrix elements?
- ▶ so far, use only a single thread (no openMP support)
 - ▶ Tensorflow sessions tend to allocate all of the available GPU memory to avoid fragmentation...

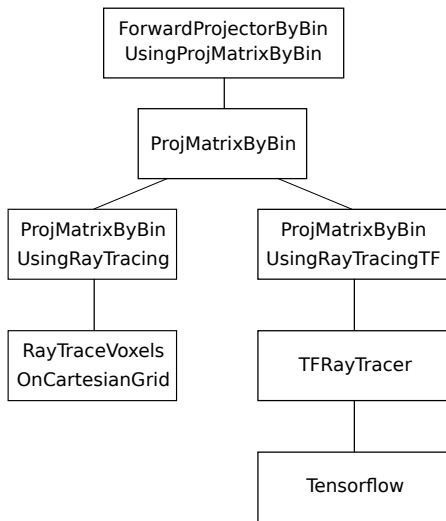
Where to find the code?

- ▶ STIR-TF: <https://github.com/philippwindischhofer/STIR/tree/stir-tf>
 - ▶ in particular:
<https://github.com/philippwindischhofer/STIR/blob/stir-tf/documentation/statistics/doc/main.pdf>
- ▶ ray tracing scripts:
<https://gitlab.phys.ethz.ch/luster/tf-raytracing>
(ask Werner for access rights)

Any comments and contributions are welcome!

Backup

Tensorflow in STIR



Tensorflow in STIR

